

Design and implementation of an application using motion input for controlling computer games as part of physical therapy

Christiaan VANBERGEN

Dimitri VARGEMIDIS

Supervisor: prof. dr. ir. L. Geurts

Co-supervisor: ir. K. Vanderloock

Master Thesis submitted to obtain the degree of
Master of Science in Engineering Technology:
Electronics-ICT, Internet computing

©Copyright KU Leuven

Without written permission of the supervisor(s) and the author(s) it is forbidden to reproduce or adapt in any form or by any means any part of this publication. Requests for obtaining the right to reproduce or utilise parts of this publication should be addressed to KU Leuven, Technology Campus Groep T Leuven, Andreas Vesaliusstraat 13, B-3000 Leuven, +32 16 30 10 30 or via e-mail fet.groupt@kuleuven.be.

A written permission of the supervisor(s) is also required to use the methods, products, schematics and programs described in this work for industrial or commercial use, and for submitting this publication in scientific contests.

Acknowledgements

Designing and implementing a user interface utilizing the relatively new Kinect camera is a challenging, albeit fascinating subject. We encountered many problems along the way, for which we enjoyed coming up with creative solutions. In the process, we were supported by many people we would like to thank.

This thesis would not have been possible to write without our supervisor professor Luc Geurts and co-supervisor Karen Vanderloock from e-Media lab, who provided us with this subject, invested their time and shared their knowledge and experience with us, allowing us to come up with solutions that have lead to the result as it is today.

During the year, we've had several occasions to discuss the subject of machine learning and support vector machine (SVM) with our fellow students Victor Martens and Alexander Kerckhofs, who helped making us aware of the pitfalls and shortcomings of the used techniques and joined us to discuss the requirements and come up with a creative solution concerning the use of SVM.

User testing is very important when designing an application with a specific audience in mind. As such, we are grateful for the useful feedback of physical therapist Dries Lamberts of Windekind, who took the time to meet us twice and evaluate both our prototype as well as the final result. But even more importantly, he shared with us his passion for taking care of children with disabilities, their need for physical exercise and the importance of making sure this is both meaningful and entertaining.

Last but not least, we would also like to thank our parents, family and friends, not only for their unwavering support during the course of past year while working on this master's thesis, but also for trying out our application and giving valuable feedback.

Abstract

Using games as part of physical therapy improves patient motivation to perform the necessary exercises. The developed application allows physical therapists to choose exercises that fit the needs of the patients and record them using the Kinect 2.0 camera. When the patient executes an exercise, machine learning algorithms are used to recognize the exercise, which activates the keyboard key linked to that exercise. This makes it possible to interact with any available game relying on key inputs.

Research is conducted on the subject of human-computer interaction (HCI) involving the use of the Kinect camera to interact with the application in what is called a natural user interface (NUI). Several prototypes are designed that focus on different interaction principles and are divided into two categories: the mime pattern and the hints pattern. The chosen implemented prototype of the interface allows users to interact with the menus and items appearing on-screen without the use of traditional buttons. Interactive elements like pulling ropes are present in the interface, providing both feedforward and feedback to the user in an effort to simplify the interaction process.

As part of the user-centered design approach, the interface is evaluated by conducting a user test with a physical therapist. Feedback is obtained both during the prototyping phase, as well as after having implemented the prototype. A conclusion from this test is that it takes more time to understand how some elements can be interacted with. With a short explanation, it is possible to decrease the time required to achieve this. The application as a whole has a low learning curve as the focus on simplicity results in the user being familiar with all required actions after recording an exercise once.

Key words: exergames, gesture recognition, Kinect 2.0, physical therapy, support vector machine, user interface

Extended abstract

Kinesitherapie is vaak essentieel bij revalidatieprocessen of om de spieren te trainen voor mensen met lichamelijke beperkingen. Computerspellen kunnen deze sessies en bijhorende oefeningen aangenamer maken om uit te voeren en betekenen een belangrijke vorm van motivatie, in het bijzonder voor kinderen. Het probleem met deze spellen is dat ze duur zijn om te ontwikkelen en aan te kopen en dat de oefeningen die een bepaald spel aanbiedt op voorhand vastliggen. Daarnaast is het voor kinderen eentonig om telkens hetzelfde spel te moeten spelen en verliest het spel op die manier snel hun interesse. Hierdoor zijn deze spellen niet altijd zinvol en effectief voor alle patiënten in het kader van de therapie.

De ontwikkelde applicatie biedt meer flexibiliteit en laat toe om reeds bestaande spellen te spelen door middel van bewegingen. De applicatie is gericht op kinesisten en laat hen toe om bewegingsoefeningen op te nemen die zinvol zijn voor de patiënt. Het is dan mogelijk om naar eigen keuze een knop van het toetsenbord toe te kennen aan elke opgenomen oefening. Wanneer de patiënt vervolgens deze oefening uitvoert, herkent de applicatie deze beweging en wordt de knop die eraan toegekend is virtueel ingedrukt. Zo is het mogelijk om met bewegingen eender welk spel te spelen dat gebruik maakt van toetsenbordbediening.

De kinesist kan oefeningen opnemen met behulp van de Kinect 2.0-camera. De applicatie gebruikt machine learning en support vector machine (SVM) om deze oefeningen aan te leren. Het geïmplementeerde algoritme splitst elke beweging op in een aantal kleinere deelbewegingen en classificeert deze elk onder een andere label. Tijdens het spelen herkent het algoritme een beweging als alle deelbewegingen ervan zijn waargenomen in de juiste volgorde. Deze benadering laat toe om de rekenkost van bewegingsherkenning te beperken en zorgt voor de patiënt voor een betere responsietijd na het uitvoeren van een beweging.

Om bewegingen op te nemen, maakt de kinesist gebruik van de grafische gebruikersinterface. Het is mogelijk om interactieve elementen van de initiële prototypes onder te verdelen volgens drie categorieën: mime-elementen, hints-elementen en WIMP-elementen. Mime omvat elementen die herkenbaar zijn uit het dagelijks leven, zoals handvaten en trekkoorden. Door een combinatie van herkenbaarheid en affordances heeft de gebruiker reeds een notie van hoe interactie mogelijk is door middel van deze metaforen, waardoor het leerproces minder groot is voor de gebruiker. Bijgevolg moet hij minder moeite doen om te onthouden hoe de interactie werkt. Hints-elementen verwachten handelingen van de gebruiker die niet teruggrijpen naar objecten uit het echte leven, zoals een scrollbar met bijhorende veegbewegingen om ermee te interageren. WIMP staat voor *windows*, *icons*, *menus* en *pointer*. Deze elementen berusten op klassieke interface-elementen zoals scrollbars en menu's. Na gebruikerstesten met een kinesist is een prototype weerhouden dat berust op een combinatie van zowel mime-elementen als WIMP-elementen.

Bij de uitwerking van dit prototype ligt de nadruk op een interface die de acties van de gebruiker koppelt aan de functies van de interactieve elementen. Met andere woorden reageren deze elementen op de manier die de gebruiker verwacht. Zo bewegen ze met dezelfde snelheid waarmee de gebruiker zijn hand beweegt over een element en reageren ze op interacties door bijvoorbeeld van kleur te veranderen.

Daarnaast reageert de interface, perceptueel gezien, onmiddellijk op gebruikersinteractie met grafische elementen en geeft de gebruiker hierover feedback.

Het geïmplementeerde prototype van de applicatie is geëvalueerd door middel van een gebruikerstest. De focus van deze test ligt op het gebruiken van de applicatie om bewegingen op te nemen, opnames te herbekijken en spellen te spelen, gebruik makende van opgenomen bewegingen. De test maakt duidelijk dat interactie met de verschillende grafische elementen niet altijd natuurlijk aanvoelt. Zo is het bijvoorbeeld wel duidelijk dat de gebruiker aan een koord kan trekken ter hoogte van het handvat, maar niet dat het nodig is om de hand te sluiten om zo het effect van het vastgrijpen na te bootsen. Er is echter beperkte toelichting nodig om er zelf achter te komen hoe het mogelijk is om met de grafische elementen te interageren. Nadat de gebruiker dit eenmaal zelf heeft gevonden, is de benodigde tijd voor dezelfde interactie herleid tot een minimum. Om dit initiële leerproces te versnellen, is het mogelijk om de gebruiker een korte uitleg aan te bieden over manier waarop er interactie mogelijk is met de verschillende elementen. Verder toont de kinesist interesse in het toepassen van de applicatie in de praktijk en ziet hij eveneens andere mogelijkheden die deze applicatie biedt, zoals het gebruik van gebarenherkenning bij kinderen met auditieve beperkingen of een interface voor deze applicatie die gericht is op het gebruik door kinderen.

Het resultaat van deze thesis is een applicatie die toelaat om bewegingen op te nemen en deze te gebruiken om met een bestaand programma of spel te interageren. De applicatie werkt onafhankelijk van zowel de gekozen bewegingen als de toepassing die zij aanstuurt met deze bewegingen. Dit laat de kinesist toe om voor de patiënt relevante bewegingen op te nemen. Opnemen gebeurt via interacties met grafische elementen die middels feedforward communiceren hoe deze interactie kan plaatsvinden. Er is een leerproces tijdens de initiële interactie met de interface, maar door het gebruiken van elementen die in het dagelijks leven kunnen voorkomen is het mogelijk om na enkele pogingen zelf te achterhalen hoe deze interactie mogelijk is. Door een minimum aan verschillende interactieve elementen te gebruiken, is het mogelijk voor de gebruiker om na het opnemen van een beweging reeds vertrouwd te zijn met deze elementen en kan hij de applicatie autonoom en zonder verdere toelichting gebruiken.

Contents

1	Introduction	1
1.1	Discussion of the problem	1
1.2	Purpose of the research	2
2	Related work	3
2.1	Human-computer interaction	3
2.2	Exergames as part of physical therapy	4
3	Design	5
3.1	Description of the application	5
3.1.1	Principle	5
3.1.2	Setup	6
3.2	Graphical user interface	8
3.2.1	Hierarchical Task Analysis (HTA)	8
3.2.2	Brainstorm experiment	9
3.2.3	Prototypes	10
3.2.4	Prototype user test	16
3.2.5	Description of the interface	17
3.3	Gesture recognition	23
3.3.1	Support vector machine	23
3.3.2	Approach	24
4	Implementation	27
4.1	Graphical user interface	27
4.1.1	Class diagram	28
4.1.2	Flow of the program	30
4.2	Back-end software	30
4.2.1	Class diagram	30
4.2.2	Flow of the program	32
4.3	Gesture recognition	32
4.3.1	Recording a gesture	32
4.3.2	Predicting a gesture	33
5	Results	35
5.1	Usability test	35

5.1.1	Description of the test	35
5.1.2	Evaluation	37
5.1.3	Interpretation	38
5.2	Technical evaluation	39
6	Discussion	41
6.1	Reflection on the results	41
6.2	Reflection on the process	42
6.3	Future work	43
7	Conclusion	45

List of Figures

3.1	<i>Overview of the interaction between a person and a computer game under conventional use</i>	5
3.2	<i>Overview of interaction between a person and a computer game using the developed application</i>	6
3.3	<i>The Kinect 2.0 camera connected to the adapter</i>	7
3.4	<i>Setup for the Kinect camera</i>	7
3.5	<i>The first prototype screen</i>	8
3.6	<i>The simplified HTA plan</i>	9
3.7	<i>(top-left) a: the UI with a real image of the user, (top-right) b: the UI with a semi-transparent shadow, (bottom-left) c: the UI with a puppet, (bottom-right) d: the UI with only hands visible.</i>	11
3.8	<i>The first prototype screen</i>	13
3.9	<i>The standard screen of the third paper prototype</i>	14
3.10	<i>The record screen of the third paper prototype</i>	14
3.11	<i>The first tutorial screen of the last paper prototype</i>	15
3.12	<i>The second tutorial screen of the last paper prototype</i>	15
3.13	<i>The third tutorial screen of the last paper prototype</i>	16
3.14	<i>An example of how the user would interact with and element of the last paper prototype . .</i>	16
3.15	<i>The regular view of the end design of the UI</i>	18
3.16	<i>Reaction of the rope handle: A) when hovered over, B) when grabbed</i>	19
3.17	<i>First view of the recording screen</i>	20
3.18	<i>View when the program starts recording, between this and the first view there is a count-down from 3, after the recording is started the user must stand still for a certain amount of time to stop, after which the text in the window will read "Done!"</i>	20
3.19	<i>The recorded element six is replayed to the user by hovering over it with any hand.</i>	21
3.20	<i>Reaction of the scrollbar: A) when hovered over, B) when scrolling is activated</i>	22
3.21	<i>The delete sequence: A) the user in the process of deleting, the red progress bar indicates how far until deletion the user has to go , B) The endpoint of the action is reached, the item turns red and fades after which it is deleted</i>	23
3.22	<i>An executed gesture, split up into four parts</i>	25
4.1	<i>The full screen of the program</i>	27
4.2	<i>The class diagram of the application, focusing on the GUI</i>	28
4.3	<i>The class diagram of the application, focusing on the back-end</i>	31
5.1	<i>Screenshot from Space Invaders</i>	36

5.2 Screenshot from Sokoban Geek	36
--	----

Chapter 1

Introduction

There are many different reasons for requiring physical therapy. They vary from recovering from a car crash to being born with physical disabilities. The goal of these therapy sessions is respectively to make sure that the patient can recover as much as possible from his injuries or to train the muscles, preventing their situation from deteriorating.

Especially for children, the physical exercises performed during therapy can be demotivating. Combining these exercises with playing computer games leads to an increased willingness to continue with the exercises [1]. However, the classic approach related to games that incorporate physical exercises is not economically sustainable.

The purpose of this thesis is to present a more sustainable solution to this problem in a way that offers more flexibility to the therapist, both in terms of the exercises that need to be done by the patients and the games they can play by performing the exercises.

1.1 Discussion of the problem

Patients often see exercises as a part of physical therapy as being fatiguing, monotonous and tedious. This problem is even more prominent with young patients and can quickly demotivate them, especially when the exercises are uncomfortable or painful to perform.

Computer games already exist that can support physical therapies. For instance, there are games that run on Microsoft's Xbox console and accept user input through the Kinect 3D camera, or Nintendo's Wii console that use a controller with accelerometers and gyrosensors. Additionally, there are games that can run on a computer, using any combination of sensors for user input, and are specifically made for use by physical therapists and their patients. However, all of these games have in common that they are very static by nature and are not often suited for the specific needs of a patient. For instance, a game that focuses on arm movement is not very effective at exercising the leg muscles of a patient. In other words, concerning the therapy, this game is meaningless for a patient that had leg surgery.

In addition, these static games have fixed input gestures patients have to perform in order to control the action on screen. Even if the gestures perfectly fit the exercise requirements for a specific patient, the decrease in attractiveness of playing the same game over and over again is problematic and loses its long-term effect. Physical therapist Dries Lamberts has experience with children with disabilities and uses computer games as a form of exercise. The initial reaction of the children to these games is positive. They are more engaged in doing physical exercises and feel motivated while doing so. On the downside, this effect only lasts until the novelty of the game wears off. After that, the children lose interest in the

game and, as such, in performing the exercises.

As games are expensive to develop, and motion-based games in particular, there is no wide variety of games that are tailor-made for people with specific needs and at the same time offer varied gameplay mechanics to remain interesting over longer periods of time. On the other hand, producing these kind of games is not economically sustainable, as from the patient's point of view, the right type of exercise needs to be supported by the game and the game itself has to be considered fun as well by the patients for them to keep invested in it.

1.2 Purpose of the research

This thesis focuses on an application that allows for a more flexible and dynamic solution to the above discussed problem. It lets the physical therapist choose what exercises the patient has to do and how these can be used as an input to control any computer game of choice. All of this can be done without requiring the therapist to have any programming knowledge.

The goal is to have an application the therapist can control mainly using the Kinect camera. It needs to feel simple, efficient and natural to do, in order to minimize the setup time before a patient can start playing. It is necessary to research the type of interface that is required to meet these objectives, in addition to finding out the easiest way to interact with an on-screen application using a camera for input.

Chapter 2

Related work

The literature study focuses on the aspects of human-computer interaction (HCI) related to gesture-based input and natural interfaces. This is followed by a closer look on research done that is directly related to using exergames as part of physical therapy.

2.1 Human-computer interaction

Gesture-based interfaces gain popularity due to the advancing technology of sensors and processors [2]. Because of this, technologies like the Kinect camera become more affordable and accessible for individuals. While interfaces supporting gesture input are referred to as natural user interfaces (NUI), this is more of a marketing term than an actual scientific designation. Natural interfaces don't feel natural to interact with just because they rely on gesture-based input. The user does not really hold the object shown on the screen to be able to interact with it naturally, so the feedback a user can get is limited [3]. Direct manipulation of objects that interact with an application offer more natural feedback [4]. The downside is that additional objects are required for interacting with the application.

Notwithstanding the limited feedback when not using additional objects, gestures are powerful tools for human-computer interaction. However, gestures as an input offer a wide variety of options and with increasing popularity, a wider variation of interfaces and gestures to interact with them emerge. This is problematic as users need to get familiar with a gesture-based interface each time they encounter a new one. The full potential of this kind of interfaces can only be achieved when some kind of standard is developed and widely used [3]. Frameworks based on reality-based interaction can be used to unify different research areas concerning interface design and to compare these designs to one another [2].

When designing an interface and ways to interact with them, it is essential to make clear what actions are required through feedforward and what the results of those actions are through feedback. Because of that, feedback and feedforward are powerful tools. Designing a gesture-based interface comes down to trying to have natural coupling between action and reaction. Natural coupling can be evaluated by six aspects: time, location, direction, dynamics, modality and expression.[5]

The time aspect requires that there is no perceivable delay between the execution of an action and the result shown on the screen. The direction and location of action and reaction have to match as well. The modality is tightly coupled to the dynamics and direction, as it must be possible to visually perceive the result of an action and the visual representation of the reaction must be similar to the action in terms of dynamics, meaning that the speed and acceleration of both have to match. The aspect of expression is also related to the dynamics of the action. The interface has to visually represent to way the user performed an action. If the action is done quickly and not very accurately, it must be reflected by the interface.

If feedback as well as feedforward are able to support these six aspects, the user perceives the interface as more intuitive to interact with and action and function become more closely related. The more aspects are fulfilled, the higher the natural coupling is. [5]

2.2 Exergames as part of physical therapy

Using games as part of physical therapy has been the subject of many research papers [6][7][8]. Since it is found that exergames improve motivation [1] and offer physical, social and cognitive benefits, they are more frequently used by physical therapists [9][10].

While the reliance of the patient on the therapist can be reduced by using exergames, the rate of rehabilitation does not increase [11]. Because of a decreased reliance on the therapist, he can focus on supporting the patient during play [12]. Since these games can be a tool of aid to the therapists, it should not hinder them in any way, nor increase the amount of work or effort they put into the preparation of a therapeutic session. For this reason, setting up everything needed before starting to play should not take up too much time. This includes but is not limited to: the time to set up the used sensors, the time the application takes to start up and the time is needed for the application to process data.

Within the same mindset of make live easier for physical therapists, the application can be a useful source of information for the therapist, as it can provide him with details of the progress his patient is making either during one session or over the course of long-term therapy. Again, this allows the therapist to focus on the essence of the therapy instead of on taking notes, keeping close track of how the patient is doing [12].

Chapter 3

Design

The focus of the application is to simplify the setup process for the therapist, ensuring that he does not require programming experience or knowledge about the underlying system, while still providing all tools needed for the patients to play a game using gesture input.

The graphical user interface (GUI) allows the therapist to interact with the application, providing him with feedback and feedforward on inputting exercises for the patients. Gesture recognition is done by means of machine learning using support vector machines (SVM).

3.1 Description of the application

The application uses Microsoft's Kinect 2.0 camera as an input device. The principle is that games can be played by first recording gestures and then performing them. In order to set up the application, there are a number of steps to be taken and guidelines to keep in mind.

3.1.1 Principle

When playing a computer game the conventional way, a person can interact with the game by pressing a keyboard key. Pressing the key then results in an action on the screen (see figure 3.1).



Figure 3.1: Overview of the interaction between a person and a computer game under conventional use

The developed application can be seen as an additional element between a person and pressing a keyboard key for controlling a game. In short, the application allows a person to remotely press a key by performing an exercise chosen by the physical therapist (see figure 3.2).




Figure 3.2: Overview of interaction between a person and a computer game using the developed application

More in detail, the physical therapist comes up with exercises that fit the needs of a patient. He uses the Kinect camera to interact with the application via a GUI. Next, the therapist uses the application to record what exercises need to be done by the patient. Using all of the recorded exercises, an SVM model is created, which is used to predict what exercise is performed by the patient while playing a game.

The therapist assigns a keyboard button to each of the exercises. This means that when the patient mimics one of the therapist's exercises, a keyboard button is pressed. To start playing a game, the therapist can choose any game or let the patient choose his preferred game. This can be any type of computer game and includes, but is not limited to: browser games, games that need to be downloaded and installed, pre-installed games that come with the operating system,... If this application is running in the background while a computer game is opened, performing an exercise indirectly presses the button that is linked to the exercise. By doing this, the patient can interact with the game.

By mapping exercises to keyboard buttons, a vast amount of existing games can be played using gesture-based input. It is however limited to button presses. Pointing with the cursor like when using a mouse is not supported by the application. The reason is that every gesture performed is seen as one single action. The gesture of pointing to a specific spot on the screen could overlap with one of the exercises the therapist wants to use. This results in undesirable behavior and negatively impacts the gaming experience for the patient. A way to solve this problem is to predefine a specific gesture that activates the *pointing mode*. In this mode, it is only possible to point to something on the screen or switch back to the *gesture mode*. However, this gives more responsibility to the therapist, who has to remember what that specific mode changing gesture is and that he cannot input that gesture as an exercise for the patient. In addition, defining a preset gesture means that not every patient has the ability to perform it, as some may not be able to move an arm or a leg, for instance. Since, from an end-user point-of-view, this setup is confusing and undoes the application's elegance of simplicity, only games can be played that require only button presses as an input.

3.1.2 Setup

Before the application can be used, the user has to download and install the necessary drivers in order to use the Kinect 2.0 camera. This is only required once per device it is used on. To connect the Kinect 2.0 camera to a computer, an additional adapter is needed so it can be connected using a USB 3.0 port (see figure 3.3). The use of the Kinect camera requires a socket. For using the Kinect, it is necessary to install a driver. This comes as part of the *Kinect for Windows SDK 2.0*. Microsoft's download page [13] states the system requirements for using the Kinect 2.0 camera. These requirements include: a 64-bit processor, dual-core 3.1 GHz or faster processor, USB 3.0, 4 GB of RAM or more, graphics card that supports DirectX 11, Windows 8 or 8.1, Windows Embedded 8 or Windows 10.




Figure 3.3: *The Kinect 2.0 camera connected to the adapter*

The camera has to be placed horizontally in order to function correctly. The distance between the user and the camera may vary between 1.5 and 4 meters, depending on the height of the user and the size of the used screen (see figure 3.4). The camera can be tilted upwards or downwards. The chosen angle depends on the height of the camera. Both can be chosen freely by the user, but the user has to ensure that he is fully visible at the position he stands. The GUI allows him to verify this. A possible configuration that works well is to have the Kinect camera pointing straight ahead on a table with a height of approximately 0.8 meters.




Figure 3.4: *Setup for the Kinect camera*

Before starting the application through means of an executable file, it is recommended to put this file in a directory that does not need to change over the course of multiple uses. The reason is that a subdirectory is created in the same directory, containing all data files of saved exercises. As long as this data folder and the executable are in the same directory, all previously recorded exercises are loaded when starting the application. If they get separated, the application still can be used, but it appears as if there are no saved exercises and it creates a new data folder in the directory the executable is stored.




Figure 3.5: The first prototype screen

3.2 Graphical user interface

A large part of this thesis was focused on the design and implementation of a user-friendly interface that can be controlled via input from the Kinect 3D camera. The reasoning behind this is that the user does not want to run back and forth from the standing position in front of the Kinect and the keyboard. Many interfaces have been design for use with the Kinect but they are generally only used to get to the game and are therefore made up of standard WIMP elements (Window, icon, menu, pointer) such as button and scrollbars to select your option. In this section the process of developing the user interface and the properties of the coded prototype are discussed.

The goal in this part of the thesis was to find a new way of interacting with a user-interface using the Kinect camera that feels natural to the average user, this despite the claims by Norman [3] that natural user-interfaces (NUI), the term that is often used to describe interfaces controlled with 3D cameras, are not natural because there is not the same feedback as when acting on objects in real life. Standard WIMP elements definitely do not feel natural in a NUI, pointing and clicking on buttons with a pointer may have already become second nature to modern humans, but it is a whole other thing when they are large and floating in the sky or fixed to a wall and the user has to slap them or close their fist over them without getting any inherent feedback as Wensveen [5] called it. For this reason element such as push buttons, scrollbars and floating windows were avoided. Additional hardware was not desired and voice control was also not considered because it is an area where plenty of research has been conducted already and implementing such an system is impossible considering the limited time for this project. If possible we wanted to find new idea's on how to control a NUI, if that proves impossible, we at least wanted to improve upon existing concepts.

3.2.1 Hierarchical Task Analysis (HTA)

To figure out what steps the user needs to perform to use the program, a Hierarchical Task Analysis was performed. This is a process in which all necessary tasks and subtasks were put into order to get a sense of the properties and elements that the UI must possess. This started as a very technical plan of the tasks that were needed but was quickly distilled into a simple visual plan that can be shown to the user to give them a good idea of what is required of them. This plan can be seen in figure 3.6 .




Figure 3.6: *The simplified HTA plan*

3.2.2 Brainstorm experiment

A brainstorm session with the eventual user, the physical therapist, would give us fresh ideas on how the user will expect to interact with a user interface using his body or sound. The expectation was that two distinct observations would be made: how the user would wish to interact with the user interface and how they would expect the flow of the program.

The idea is to keep the subject from being influenced by our ideas or by obvious solutions. To start only the basic concept of our application and the abilities of the Kinect camera are explained. Then the subject is asked how he would imagine every aspect of the application, a few questions from the script are: What information would you expect to see on the screen? How would you start a recording? What would you want to see during recording?

To find out whether such a brainstorm session would yield any viable results, brainstorm sessions with our relatives were conducted first. Two iterations were performed with four subjects, the first test was with two males and the second was with two females, of both genders there was one person around 25 years old and one around 55 years old. Each subject was interviewed completely separate to minimize the influence they had on each others ideas.

The sessions were very taxing for the test subjects, the concept we tried to introduce was too abstract and foreign to them. This causes them to be shy with proposals, even after being ensured that there are no bad ideas. It seemed to give them the feeling that they are not smart enough to make good proposals. As a result they fell back on proposing familiar UI elements or concepts such as push buttons or voice control. Getting locked into a mindset that we were designing games to play with the Kinect was also a minor issue, thereby answering the questions with that in mind, which was not the point. After two iterations with minor changes between them, we decided not to subject the physical therapist to this kind of questioning. Setting up a meeting would waste too much of his precious time with little gain to us. Worse, it might instill the physical therapist with a feeling of dread for our next meeting which might reduce his level of co-operation.

Instead the decision was made to brainstorm with each other and our supervisor professor Geurts to develop a few prototypes that we can eventually show to the physical therapist to see which one seems the most user-friendly. The two major interaction patterns that came out of this brainstorm session are, for the purpose of this thesis, named the hints and mime patterns. In the rest of this chapter the abbreviation UI is used to say user interface and a UI element is a button, scrollbar or any other element that the user can interact with. When discussing the different concepts we often refer to a movement by the user or

position that is taken to activate a function in the program as a sign.

The hints pattern implies that the screen is largely devoid of UI elements and actions are performed when the user does a certain movement or takes a certain position, for example, the user forms a cross with his arms to delete a recording. These signs have work better when the action has a similar meaning to the meaning of the sign in the users everyday life, as in the last example where an X is related to delete because it generally means bad, wrong or unwanted in Western culture, but this might be different for other cultures. The lack of feed forward that this implies means the user should learn all of the signs at the start of application, putting a large strain on the user's cognitive ability which is not ideal. Alternatively, all possible action could be displayed on the screen around the user with a depiction of the sign. Though this would put less strain on the user's cognitive abilities, he would still have to focus on these pictures and try to figure out what sign they depict. But what if the sign is a movement? It is impossible to depict it statically, so how can you make this clear without distracting the user with constant moving pictures? A more significant problem might be the fact that positions and movements might be recognized by accident, an undo movement could partly solve this but you also need a redo button to compensate for any accidental undo-activations. The only real benefits to using the hints pattern is that activating an action can be done from anywhere and it could be faster once the user has mastered the signs, though this depends on the ease of performing the signs and the delay between a sign and the corresponding action in the program.

The mime pattern, as the name suggests, is the act of mimicking an action that is needed to act on everyday objects such as opening a door. But where an mime artist does this without any indication of that object being there, here an image of ,for example, the door would be displayed on the screen. The technical term for such a UI element is a metaphor because the function is similar to what a user would do with the object in real life. The power of this pattern is that the user should immediately recognize the UI element and know how to interact with it because it has the same affordances as the object in real life. To achieve the best effect, the UI elements that are chosen should need signs that the user has done frequently already, so that doing them feels natural and familiar. Here it is less important than in the hints pattern that the sign performs a function that is similar the real life such as going to a different screen when opening the door, but it does improve it. Wensveen [5] already said that for this to work correctly feedback is crucial, the reaction of the UI element must match the expectations of the user otherwise he will be confused and the flow of the program is interrupted. Disadvantages are that the user cannot activate every action from anywhere and it might be slower because of that.

Our opinion is that one should not mix the two patterns because the user will expect to interact with everything on the screen when using the mime pattern, if one action is activated using the hints pattern without indication the user will be confused when he has forgotten the sign. Even if there is indication the user will likely try to interact with the indicator of the sign as a button which will not have the desired effect. Combination with regular WIMP UI elements, such as push button or scrollbars, is possible for both patterns but it combines much better with the mime pattern since the user is already interacting with elements on the UI while pointing at a specific area in the UI doesn't fit well in the concept of the hints pattern and would even detract from the idea of being able to do every action everywhere. To sum up, in the mime pattern the behavior of the UI elements are simple and obvious but possibly more inefficient, like pulling a handle on a slot machine to spin the slots. While the hints pattern is more abstract and complex, but should allow a more efficient interaction once mastered.

3.2.3 Prototypes

With these patterns in mind the design of the prototypes focuses more on designing UI elements and methods with which the user can interact that are simple, fast, ergonomic, feels natural and is not taxing to the user physically or cognitively while being practical and not easy to activate accidentally. The regular way of testing a concept UI is by doing paper prototyping which means that the user is presented with a paper version of the UI (hence the name), whenever the user touches a UI element the designer manually

changes the paper prototype to reflect the activated action. This kind of testing is absolutely not suited for the goal of this project, it is more to assess the placement and shape of UI elements and the general flow of a program. What is needed here is a way to interact with the prototypes in a way similar to how the real program would work but without having to code the whole program. For that reason we took two example programs from Microsoft's Kinect SDK and made some adjustment to show the user in front of the prototype.

When designing a motion controlled UI it is very important to let the user know what the camera sees of their signs. The manner in which this feedback is given will have a great impact on how natural the user's experience with the UI will be. There are four viable options:

1. A real image of the person as the camera films it. (**Figure 3.7 a**)
2. A semi-transparent shadow of the user. (**Figure 3.7 b**)
3. A puppet that copies the user's movement. (**Figure 3.7 c**)
4. No body is shown, only images that portray the position and state of the hands. (**Figure 3.7 d**)




Figure 3.7: (top-left) a: the UI with a real image of the user, (top-right) b: the UI with a semi-transparent shadow, (bottom-left) c: the UI with a puppet, (bottom-right) d: the UI with only hands visible.

A real image of the person is the most identifiable for the user because it is literally an image of himself, but this doesn't mean that it would be the best option to use in interacting with the UI. Though it has the

advantage that it is crystal clear which movement the user carried out, which is good for recording and replaying exercises. The most important issue is that the image is taken from the front so the image of the screen is facing away, so unless the UI elements are drawn on top of the user's image it is hard to imagine you are grabbing a handle or pushing a button on the UI. Placing the Kinect camera behind the user to solve this problem, is technically not possible because an important part of the Kinect's person recognition is based on recognition of the facial features. This would also require a very large amount of room to operate because a distance of at least a one meter from the screen in front of the user and at least 1.5 m from the Kinect at the back needs to be maintained. When the UI elements are drawn behind the user, he will also obscure any elements directly behind him. Another issue is that most people are not comfortable with an image of themselves being displayed on the screen which could detract from the user's experience. The physical therapist we interviewed, who works in a school for disabled children, also noted that this might cause legal problems with parents who do not want unauthorized footage of their child being stored. A final technical issue is that displaying the real image of the user is very resource intensive and in this implementation it would sometimes cause the program to lag.

A semi-transparent shadow would solve the legal issue, the discomfort experienced by users and it also solves the problem of the image facing away from the UI because there are almost no features that indicate which way the image is facing, making it look like the shadow is facing towards the UI. By making the image semi-transparent any UI element behind the user can still be seen. A disadvantage is that it is harder to identify subtle differences between similar recordings. It is also still very resource intensive, trying to optimize this was beyond the scope of this thesis.

A puppet that copies the movements of the user has similar benefits to the shadow and can be transparent or non-transparent. It is far less resource intensive but makes it even harder to distinguish similar recordings from each other.

Showing no body and only hands can obviously only be used to interact with the UI and not for the recording of an exercise. It removes all distractions from the screen and shows only the information that is essential for the user to interact with the UI. But because of this the interaction might feel less natural to the user.

The shadow type of representation has the most advantages so it is chosen to make the most prototypes with. Though in a experiment to see if the user is not bothered with using more abstract representations to increase efficiency, the representation that only shows images of the hands of the user is also used in the prototypes.

In total six prototypes are created ranging from obvious UIs using mostly mime pattern elements and metaphors to more abstract UIs where more elements of the hints pattern and standard WIMP concepts are used with less mime pattern elements. The first four prototypes are only the screen in which the user can record one exercise multiple times to provide the data for the SVM (step 2 to 4 in figure 3.6). They use a shadow image of the user to control the UI. The other two prototypes are menus to browse to a recording and replay it while the user only sees images of his hands. The following list sums up all the prototypes:

1. A UI with mostly mime pattern elements and metaphors.
2. A UI that combines mime patterns with standard WIMP elements.
3. A UI that also combines mime patterns but with more WIMP elements.
4. An alternative mime pattern element for the scrollbar only.
5. A UI that combines the hints pattern with WIMP elements.
6. A UI that combines the hints pattern with more WIMP elements.

In this thesis only three of these prototypes are discussed to show the kind of elements that were experimented with, these are: the first prototype, the third and the last.

The first paper prototype, that can be seen in figure 3.8, mostly uses the mime pattern and metaphors, the orange color is used to indicate a possibility for an action. On the right the door on the side says exit and would swing open when the user grabs the orange handle and pulls to the left. On the top of the screen a collection of recordings can be seen pinned to a scrollbar with orange thumbnails. The arrows on the sides imply that the recordings can be moved. The user would grab the orange indicator bar to browse the recordings. To delete a recording the thumbnail of that specific recording can be grabbed and dragged to the recycle bin to release it there. The user can display the recording by dragging it to the projector on the left side of the screen where the big 1 is displayed. In the middle a camera with its lens pointed towards the user is displayed, by moving his hand forward on the “opnemen” button the user could start a recording.




Figure 3.8: The first prototype screen

The third paper prototype in the list combines mime pattern elements with more standard WIMP elements to streamline the process, it can be seen in figure 3.9. On the right of the screen is a lever that is activated by grabbing the orange part of the lever and pulling it downwards causing the user to leave this screen. The record button is a push button that the user activates by moving his hand forward over it. This action activates the record screen as seen in figure 3.10. The window shows what is going to be recorded, the red one indicates where the program will count down from three until the recording starts, allowing the user some time to get to his starting position. All other UI elements are grayed out to indicate that they are not available to interact with. Once the recording is finished the program returns to the screen in figure 3.9. On the left is a scrollbar that contains elements which represent previous recordings of which they show static representations. On the right of it, there is an orange scroll wheel with which the user can scroll through the scrollbar by simply hovering his hand over it. When the user stops scrolling the elements are locked into a position as is shown in 3.9. The user can slide the top element into the “THRASH” or “PLAY” area by moving over it to play or delete the element.




Figure 3.9: *The standard screen of the third paper prototype*




Figure 3.10: *The record screen of the third paper prototype*

The last prototype is a combination of the hints pattern where the user has to close his hand over the element with which he wants to interact to open a typical WIMP element called a context menu with which the user can interact. The user can only see images of his hands as discussed earlier in this chapter. It starts off with a small tutorial screen as seen in figure 3.11, when the user closes his hand within the orange box, the context menu as seen in figure 3.12 opens. When the user slides his closed fist towards any of the elements it reacts as seen in figure 3.13 to indicate the option is chosen. Throughout the prototype the user can interact with the chosen recording as seen in figure 3.14, in this instance he can choose between playing the recording, deleting it or canceling the current menu.




Figure 3.11: *The first tutorial screen of the last paper prototype*

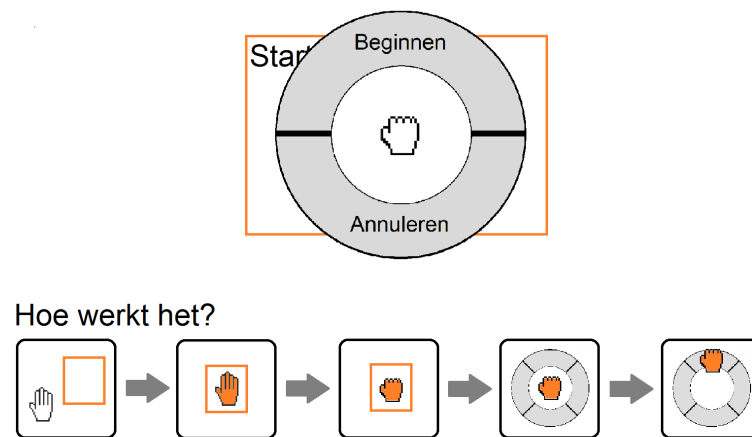


Figure 3.12: *The second tutorial screen of the last paper prototype*




Figure 3.13: The third tutorial screen of the last paper prototype




Figure 3.14: An example of how the user would interact with an element of the last paper prototype

3.2.4 Prototype user test

The main goal of this user test is to see how the user interacts with the different paper prototypes, how fast can he do the required signs and his opinion of the signs, whether they feel natural or not and which one of the paper prototypes speak to him the most.

For these tests the Kinect is placed on a table of regular height and the user is asked to stay within two to four meters distance from the camera. The screen with the UI is placed directly behind the Kinect camera.

The user is first introduced to the concept and flow of our program through the use of the HTA plan as seen in 3.6, to bring the user up to speed with what is necessary to operate the eventual program. Then a small demo program is presented showing the capabilities of the concept and the Kinect to give him an idea of the end goal and the possibilities with this concept. After some more introductory questions the user is assured that he is never at fault and that we are responsible for any misunderstandings he might

have. Before the actual tests, permission is asked to film their proceedings. During the tests the user is reminded that we like to hear his thoughts on what he sees and thinks.

Then the user is asked to interact with the prototypes using the Kinect. To mimic the functions of the UI, a Wizard of Oz method is used to change the screens whenever the user performs the required action. Some prototypes still require some explanation but generally the user is only assisted when he is completely stuck and at the end of every prototype an explanation of the intent of each UI element is given. At the end every prototype is gone over again without the Kinect to reflect about the performance of every prototype. The user is asked to select a personal favorite.

In total four candidates were subjected to these tests of which were one physical therapist and three relatives or friends but this paragraph will focus on the results with the physical therapist supplemented with the results of the others, thereby the physical therapist is referred to as *the user* and the others as *the other users*. It is important to mention that the user has previous experience with Kinect games which sometimes influenced his reaction, such as he did not know that he could move his hand forward to mimic pushing a button. For the first prototype 3.8, the user was trying to interact with UI elements that were not intended to be interacted with, he later indicated that he did not see the metaphors as they were intended, the screen was too hectic and he didn't see which ones were buttons and which were not. He also found that the required signs were too long. Most of the other users also made the same remarks about the first prototype. A rope in one of the prototypes on which the user could pull was clear to him, but he had no idea how to interact with the scroll wheel to scroll through the scrollbar. In one of the prototypes the screen changed based on where the user was positioned, this caused a great deal of confusion, the same goes for the other users. The lever on the third prototype as discussed above, to leave the screen was clear but he did not like it, the metaphor was too foreign to him. The user was quick to pick up the controls of the last two prototypes that only showed the hands, but he later indicated that the representation of the hands only didn't feel right to him and that he much preferred to see all his options on the screen. Other users indicated that they expected the UI to react as if it were a touchscreen. When asked how they preferred to see themselves on the screen, all users experienced discomfort with the real image of themselves and the therapist made the comment about legal issues with the patient as seen before in the discussion of the user's representation. They all preferred the shadow representation. Remarkable was that all the users chose the third prototype, as seen in figure 3.9 earlier on, as their favorite, mostly due to the simplicity and clearness of it.

The conclusions are that the most important characteristics are that the screen needs to be simple and straightforward with all of their options displayed on the screen permanently and modes dependent on subtle parameters, such as the user's position, need to be avoided. The preferred representation of themselves is the shadow. For the purpose of simplifying the UI regular WIMP elements are still allowed to be used without interfering with the user experience, though clearer affordances for elements that do not look like any object in real life are needed.

3.2.5 Description of the interface

The end design looks as seen in figure 3.15. Starting from the left, the user can see a gray window in which a recording can be replayed, this is a typical WIMP element because it does not have any elements to it that can be seen as a metaphor to anything the user is familiar with in real life. Right of this, there is the rope as seen in prototype two in the **appendix**, which is a very distinct mime pattern element. Further right is a simplified version of the scrollbar containing the recordings with the currently selected recording highlighted in orange. The user can slide this element to the "DEL" bar to delete the element. The black bordered box at the feet of the puppet represents the ground on which the puppet walks so it does not appear to float in mid air.

Though we had previously established that the shadow was the most natural way of representing a user

in the interface, performance issues forced us to resort to using the puppet image. The reason why the puppet is not semi-transparent is because then the joints are more visible than the rest because two semi-transparent images are drawn on top of each other which combine into a less transparent color. This breaks the unison of the puppet's coloring making it seem less human which would be a distraction that detracts from the experience of the user. Another detail of the puppet is that the centerpoint of the spine is fixed in the vertical direction but not in the horizontal direction. This eliminates any dependency on the position of the camera so that, as long as the user is within the field of vision of the kinect, he can always reach every UI element on the screen. This measure introduces two problems. The biggest problem is that when the user tries to squat the puppet will lift its feet instead of lowering its torso, this might cause confusion and does not feel natural. The other problem is that the user cannot reach any UI element that is placed where it could only be reached by bending downwards. This is not seen as a big problem because such a sign is very taxing for an adult to perform and thus not wise to use in a UI. The puppet is free to move in the horizontal direction but there is an offset on where the puppet is drawn so that the puppet appears to stand between the rope and the scrollbar when the user is standing right in front of the kinect. In a normal setup the screen will stand behind the kinect so that the user can look straight forward when interacting with the UI. The depth coordinates in 3D space are also set to a fixed distance before conversion to a 2D screen to ensure that the puppet is always drawn to the same size no matter how far the user is standing within the field of vision.




Figure 3.15: *The regular view of the end design of the UI*

An important aspect of the UI is the color coding that is consistent throughout the UI except on the puppet itself though it should be apparent that the puppet is not much more than a representation of the user. This should give the user extra feed forward about what the consequence of their action will be. Orange is indication of interaction possibilities, green means a hand is hovering over the element, blue gives information about what happened or will happen after an action and red indicates a more serious action is about to be activated. The exact meaning of all these color codes will become clear during the description of the UI.

Since it is important for the user to know when the system sees that it is grabbing the rope, the UI provides feedback to the user by coloring the hands green when they are open and red when the user makes a fist. In figure 3.16 A can be seen that when the user hovers his hand over the orange handle of the rope the handle turns green indicating that it can now be grabbed. When the user closes his hand, the handle turns red to show the user that he is pulling it is shown in figure 3.16 B. The dynamics for pulling the rope are simple, between the starting position and the end position, which is below the starting point, the rope will make the same relative movements in the vertical direction as the user makes with the center of his

hand. In this way the interaction is linked in time, speed, direction, dynamics and location of the user's hand. When the rope is near to his end position the record screen is activated. There is no progression bar to indicate how far the user is from activating the function because generally the action of pulling a rope is very sudden and fast, as long as the function is activated before the user runs out of momentum this kind of feedback is not needed here. In the normal state the handle guards are not colored because the affordance is clear enough and a person would not normally grab a handle by the handle guards, they are however colored green and red to make it more apparent to the user that they are acting on it.




Figure 3.16: Reaction of the rope handle: A) when hovered over, B) when grabbed

The rope activates the record screen seen in figure 3.17. In this screen a separate window is opened with a scaled version of the puppet, another typical WIMP element. The large change in the UI clearly indicates that the user is in a different mode. Here the puppet is also fixed in the horizontal direction to indicate correctly how the program will see the exercise. To avoid confusion the window is opened on top of the last position in which the user was before pulling the rope. Since the goal of this screen is to record one exercise multiple times two static puppets are drawn behind the user, the light blue puppet shows the starting position and the dark blue puppet shows the end position of the exercise. Care has been taken to choose a color that does not conflict with the green of the user's puppet and doesn't overpower it by being too bright while still being clearly visible and distinguishable. The colors also follow the color code defined earlier in this chapter since it gives information about the first recording. The message "Get Ready" is shown for 1 second before changing to a count down of "3","2","1","Go!". At "Go!" the window's edges turn red to clearly indicate that the program has started recording the user's movements. An example of how this would look can be seen in figure 3.18. The program stops recording when the user remains still for certain amount of time. After that the text changes to "Done!" and after half a second the screen returns to figure 3.15. The result is that a new element in the scrollbar is added to the top. It is given the next number and set into the orange selection area.




Figure 3.17: First view of the recording screen




Figure 3.18: View when the program starts recording, between this and the first view there is a countdown from 3, after the recording is started the user must stand still for a certain amount of time to stop, after which the text in the window will read "Done!"

As said before the scrollbar contains elements that are each linked to a recording of the current exercise. In these boxes a small preview can be seen of the start and end position of the linked recording in the same color scheme as during recording. The numbering of the boxes goes from the first recording as one counting up to the last recording, though these numbers are not the same as the ID's of the gesture class as seen further on during the discussion of the implementation of the back end, the order of the numbers is still the same. When the user hovers one of his hands over the recorded element within the orange selection box the effects as seen in 3.19 are triggered. The specific recorded element is replayed within the previously grayed out window on the left. The blue bar on the bottom of the window indicates the progression of the recording. As long as the user hovers over the element in the scrollbar, the recording is repeated with half a second delays between repeats. This was a deliberate choice to make sure the user is fully aware of which recording is playing in the window because he should notice that the recording is only playing when he touches it. For the same reasons the number of the recording is also displayed on the left of the screen where it does not obstruct the user in reviewing his recording but is there when the

user is uncertain of which recording is playing. Additionally both the orange selection box in the scrollbar and the edges of the display screen are colored green, the color previously linked with hover over information. The redundancy in linking the replay with the recording comes from the fact that the screen is so far from the actual element it is connected too. The scrollbar is on the right because right handed users are more common, so it is easier for them to operate the scrollbar. The delete action is already on the right of it and can hardly be moved to somewhere where it is still as easy to operate. Following both the proximity and similarity law of the Pragnanz law's would make the most sense but placing the screen on the right of the delete button would still have the delete button between it and the recorded element and would place the user's puppet too far from the center of the screen, out of the focus of the user.




Figure 3.19: *The recorded element six is replayed to the user by hovering over it with any hand.*

To activate the scroll function on the scrollbar, the user has to hover his hand within the boundaries of the black bordered box but not within the orange selection box, the black borders of the scrollbar turn green to indicate that the user is hovering over the scrollbar, this is shown in figure 3.20 A. Scrolling only starts when the user has moved a small distance up or down to avoid to scrollbar being immediately stuck to the user which would be annoying. Once scrolling is started it continues as long as the user keeps moving in the same direction, this also counts when going through the selection box area and even when going outside of the scrollbar's borders. This makes sure that the user is not surprised by a sudden stop of the scrolling action when accidentally going inside of the selection box or outside of the scrollbar. The whole reason why the user cannot initiate a scroll within the selection box is because the element that is placed there has to be able to be hovered over to display it, as explained earlier, and it has to be able to be moved to the side for deletion. Accidental activation of the scroll function is minimized by this measure. Though there is more to this, informal tests showed that users who are unfamiliar with the UI often try to scroll down or upward starting from the selection box, the large delay in reaction confuses them. To counter this the activation area of the scroll function is expanded into the orange area for about 15 percent of the selection box height on both the top and the bottom. Once activated the elements stay on a fixed distance from each other but all elements are moved with the same dynamics as used in the record rope, they perform the same vertical relative movement as the hand that is hovering over it. During scrolling all recorded elements are locked to any interactions and the element that will land in the selection box is indicated by a light blue coloring of the white background. The example of scrolling can be seen in figure 3.20 B. When the user stops scrolling the recorded elements are locked into positions as seen in 3.20 A. Only the element in the selection box can be interacted with.




Figure 3.20: *Reaction of the scrollbar: A) when hovered over, B) when scrolling is activated*

The delete sequence of the recorded element in the selection box can be started by hovering over the middle of it. After this the element will follow the relative horizontal movement of the interacting hand between start point and end point with the same dynamics as the scrollbar and the rope. As seen in figure 3.21 A, The "DEL" box is filled up with a red progress bar to indicate how far the user still has to move until the item is deleted. When the end point is reached the recorded element turns red and fades away within a half second to clearly indicate that this recording is deleted. Once the endpoint is reached the action is irreversible and for the duration of the fading the scrollbar and all the elements are locked for any further interaction. An example of this is shown in figure 3.21 B. The reason why this action is only started from the middle is to increase the difference between initiating the delete sequence and replaying the recording and also heighten the threshold to start such a serious action as deleting an item. After deleting an element the other recorded elements are renumbered.




Figure 3.21: *The delete sequence: A) the user in the process of deleting, the red progress bar indicates how far until deletion the user has to go , B) The endpoint of the action is reached, the item turns red and fades after which it is deleted*

3.3 Gesture recognition

In order to provide enough flexibility concerning the type of exercises, SVM is used. SVM supports supervised machine learning and its use encompasses two modes: train and predict. After a discussion of the way SVM can be useful for this application, the approach relying on SVM to predict the executed gesture is explained.

3.3.1 Support vector machine

A model can be trained with a given set of data and a label to classify the gesture. The amount of numbers in one data set is referred to as the number of features. If it contains n features, the entire data set can be seen as a single point in a n -dimensional space. It is possible to train the same gesture more than once. In that case, the same label is used to indicate that the given data set is related to the same gesture. In other words, all trainings of the same gesture are linked to the same label. All of these trained gestures with the same label are seemingly similar, but actually contain variations due to noise during the measurement of the gesture or a slightly varying execution of the gesture.

After training all required gestures, a SVM model is created. Given the data set of a gesture, this model can predict which of the trained gestures has the biggest resemblance to the given data set. As a result, the label of the most similar trained gesture is returned. This also explains the necessity of having multiple trainings recorded for each gesture. Errors in a single training due to noisy measurements of the Kinect camera can lead to wrong predictions. Having multiple trainings minimizes the influence noise has on the prediction and keeps into account that the user executes all gestures with slight variations. As a result, the model can predict more accurately which gesture is performed.

However, there are some things to keep in mind when applying this strategy. Firstly, while inputting multiple repeats of the same gesture helps with predicting the gesture after a model is created, it takes more

time for the therapist to do this. Secondly, when trying to predict a gesture, the generated SVM model always returns the label of the most similar gesture, even if they are not related at all.

These problems are tackled as part of the approach to using SVM to predict gestures.

3.3.2 Approach

A gesture consists of multiple frames. Each frame contains 25 joints and each joint consists of an x , y and z component. This results in a total of 75 features that are being considered for each frame.

Two approaches are considered when it comes down to learning how to recognize gestures. By directly comparing these approaches, it is easier to identify their advantages and disadvantages.

The first approach is to add a time stamp to each frame as an extra feature, which indicates the time relative to the first frame of the gesture. This results in having 76 features per frame. The additional time stamp feature can be used to make sure that the gesture isn't just similar in space, but also in time. If performing a certain gesture takes t seconds and is captured at a rate of f frames per second, this amounts to $76 \cdot t \cdot f$ features. For a 3-second gesture recorded at 30 frames per second, which is the highest sampling speed of the Kinect camera, this amounts to 6840 features for a single training of the gesture. In other words, the number of features of one training depends on the duration of the gesture. The entire gesture is classified as a single gesture. During prediction, a gesture with the same number of features can be input to verify if that gesture matches any of the trained gestures.

There are several problems with this first approach. If the number of features of the gesture used for predicting does not match the number of features of the training gestures, the prediction is not accurate and should be discarded. Discarding is necessary as SVM always returns the label of a gesture, even if the predicted gesture is completely unrelated to any of the trained gestures. This poses a problem as different gestures can have a different duration. Even different trainings of the same gesture can take for instance 3 seconds and 3.1 seconds, implying that the recordings have a different number of frames and thus a different number of features. A possible solution is to recalculate the entire gesture and use interpolation to convert the set of frames to a new set with a known, fixed size and preferably with equidistant time stamps.

Furthermore, not just the trainings of one gesture, but all of the gestures need to have the same number of frames. The gesture to predict is not known beforehand and can only be predicted accurately if the number of features for both the predicted gesture and trained gesture are equal. As a result, it is required that all gestures have an equal amount of features. This means that short gestures need to be mathematically extended with additional frames to match the size of longer gestures. Another side effect is that all gestures need to have the same duration, not just the same amount of features. This limits the therapist in choosing exactly what exercises he wants the patient to perform. Also, predicting a gesture can only be as fast as the trained gesture with the longest duration. More in particular, assume the longest gesture is 5 seconds in length. It then follows that it takes about 5 seconds to predict any gesture. This also means that patients that control the game can only perform one action every 5 seconds. As the games to be played are unknown beforehand, a slow reaction time of the application can render some games unplayable. As such, this approach is not viable.

The second approach is to split up one gesture into n smaller gestures and classify each of them differently, assigning a different label to each part of the gesture. Assume that a gesture is split up into 4 smaller parts so that each part contains approximately an equal amount of frames (see figure 3.22). Consider a 4-second gesture sampled at 30 frames per second. The complete gesture consists of 120 frames. By splitting it up into 4 parts, each part contains 30 frames. All first 30 frames are labeled with

the same label, for instance 1. The next 30 frames receive a label 2, and so on. In other words, the considered gesture is split up into 4 postures with each having 30 slightly varying trainings. In contrast to the first approach, as described above, prediction does not happen for an entire gesture, but for separate postures. The condition for having executed the entire gesture is that each of the postures are predicted correctly and in the right order. To put it differently, n pictures are taken of the gesture and if at some point during prediction all pictures are executed in the right order, the application acknowledges the execution of the gesture.




Figure 3.22: *An executed gesture, split up into four parts*

The biggest advantage of this approach is the flexibility of it. Each posture corresponds to a single frame, which contains 75 features. This eliminates the need of having gestures with the same length or having to modify training data to have gestures with an equal number of features. It is even possible to not just link a gesture, but also a posture to a keyboard button, which allows the physical therapist to choose the exercises that fit the needs of a patient the best. Additionally, the application reacts immediately to an executed gesture, not after a fixed amount of time. This allows for a wider variety of games being playable using this application.

Another advantage is that it solves the issue where SVM always returns the label of a predicted gesture, even when the patient is not doing anything. A gesture is only recognized when all parts of it are executed. In other words, gestures are not recognized involuntarily and, as such, actions like button presses that influence the game are not executed by accident. This also means that no neutral or *wrong* gestures are required to link a certain gesture to no in-game action.

The disadvantage of this second approach is that there is no strict requirement for the entire gesture to be executed in about the same time as the gesture recorded during training. If a gesture is executed faster during prediction compared to during training, it is recognized as the same gesture being executed. However, it is possible to solve this timing issue to some extent when the gesture during predicting is performed slower. The gesture can be ignored if more than a certain amount of time passes. This time threshold takes the gesture with the longest duration into account and is chosen higher than this in order to allow all gestures to be executed and successfully recognized.

Chapter 4

Implementation

The application is developed in C++ using Microsoft's Visual Studio IDE and can run on computers with a Windows operating system. On an implementation level, platform-specific features are used for saving and loading data files and running the GUI. The implementation consist of three big parts: graphical user interface, back-end software and gesture recognition.

4.1 Graphical user interface

In this chapter the implementation of the UI is discussed. First the used technologies are given, then the class diagram is discussed per class and last the general flow of the program is described.

COMMENT: zou ik deze disclaimer houden? , de UML moet echt veel groter

WPF is used to easily create buttons and input fields to use as a developer's UI seen on the right side of figure 4.1. The interface for the actual users is created in the D2Direct. Because this part of the program was used to experiment with different types of feedback and UI elements it had to be made easily adaptable, for this reason it uses structures that are commonly seen as bad practice like global instances of classes and loosely defined class roles.




Figure 4.1: The full screen of the program

4.1.1 Class diagram

The program is build out of an uncopyrighted example by Microsoft called *BodyBasics-D2D* that contained all code to draw 2D skeletons of people seen by the Kinect camera on an empty background. The original program had all of its code centered into one class but has been refactored into a more defined structure as is seen in Figure 4.2. The original code has been spread over the Main, UI, Model and D2D_Graphics into a Model-View structure with the Main class initiating the resources for the Kinect, creating the UI and Model class and starting the program loop. The UI received all the elements retaining to the view and controller, while the model received all code retaining to the processing of the Kinect input. The D2D_Graphics serves as an interface between the UI classes and the D2DDirect canvas.

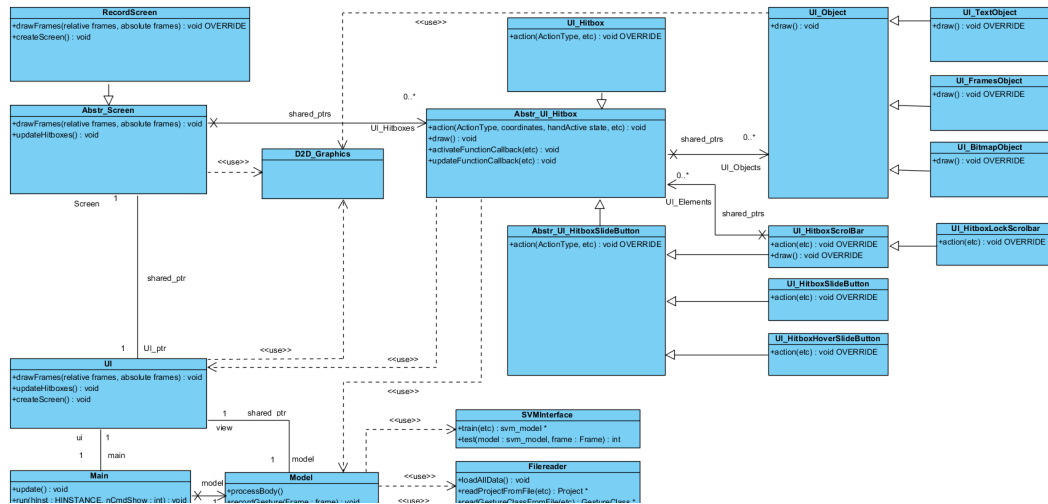


Figure 4.2: The class diagram of the application, focusing on the GUI

The UI class is the main class for the UI part of the program. It serves as the main manager of the view and controller related classes and is the only one through whom the Model class can activate UI elements. It creates the main window in which the application is displayed and handles the input from the WPF UI elements and the keyboard input. At construction a global instance of the D2D_Graphics class is created so that all other classes can have direct access to the screen. It also initializes the resources for the D2DDirect canvas. It contains a pointer to the current screen that is displayed, every time a new screen is opened the object of the old screen is freed and replaced with a new object of the current screen. This class is responsible for creating the instance whenever a new screen is opened. Though at this point there is only one screen so the pointer is never changed. Three important functions are: drawFrames(), UpdateHitboxes(), createScreen(). drawFrames() manages the D2DDirect resources and starts and ends the drawing on the canvas of D2D_Graphics and also passes the draw command to the Abstr.Screen class, it is called by the Model class to redraw the screen. UpdateHitboxes() is used by the Model whenever changes in the model data need to be displayed in the UI it is also passed on to the Abstr.Screen class. The createScreen() method is called by the Main class when all resources are initialized and the Model is created, the UI class passes it on to the Abstr_screen class to instruct it build the elements for the current screen.

The D2D_graphics class contains methods from the original example program, methods to initialise the D2DResources, to convert 3D joint coordinates to 2D screen coordinates and to draw the body, though some of these methods were adjusted to our preferences. It also contains new methods to draw Bitmaps, Text or scale 2D skeleton coordinates to different sizes.

To explain the rest of the structure it is necessary to start with the simplest class and work our way up to the bigger classes.

The UI_Object class and its children are an undisputed part of the view organized in a kind of strategy

pattern in which each child adds properties and defines the draw() function differently. It contains information about how a rectangle element is drawn using a center, width, height format whilst also containing additional properties such as color, border color and border thickness, etc. The most important function is draw() that uses the D2D_Graphics to draw their rectangle. As mention before, each child class overwrites this function to draw their own specific type of image.

The UI_TextObject class is a child of the UI_Object class that is used to draw text on the D2D_graphics. It contains only the necessary properties of the text. The rectangle within which the text is drawn is defined by the parent class.

The UI_FrameObject class is also a child of the UI_Object class that is used to draw images of puppets other than the one directly controlled by the user. The puppet is scaled to the dimension of the rectangle defined by the parent class.

The UI_BitmapObject is the last child of the UI_Object class that is used to draw bitmaps scaled to the dimensions of the rectangle defined by the parent class. This is the only object that uses states to decide which image is drawn. Three states are defined : standard, hover and handActive. Standard behavior is to draw the standard image for all states, but each state can be assigned to a different image.

The Abstr_UI_Hitbox class represents a rectangle that takes on different states dependent the position and state of the hands of the user. The different states that the class can attain are: "hover" when a hand is hovering over the borders of the defined rectangle, "handActive" when the hovering hand is in the activeHand state, such as a closed hand, "activeHandOutside" state when the user's hand moves outside of the rectangle while the hand is in the activeHand state. The states are indicated by flags in the class. The action(etc) function defines behavior for the entering, holding and leaving of a each state using enums. Additional flags allow the usage of special circumstances such as when the hand enters the rectangle while in the activeHand state. This action(etc) function is the main function that is redefined by every child to obtain more complex behavior using these states. This means that every hitbox that has a different effect on the UI must have his own class that is a child of this class. This is the reason why the strategy pattern is used. The class contains a vector of shared pointers to UI_Objects upon which the class can act, like move them or change their color and which is used to instruct the UI_Objects to draw themselves when the method draw() is called. Multiple hitboxes can have a pointer to the same UI_object so they can all change it. It also contains two callback functions, one is called the activateFunctionCallback which defines the action performed on the Model and UI when the criteria for activation of the hitbox have been reached and one the updateFunctionCallback which defines the behavior when the data in the Model is changed and the UI is updated accordingly.

The UI_Hitbox class is a child of Abstr_UI_Hitbox and as mentioned before it redefines the action method to define a specific behavior when entering, holding and leaving each state. This type of hitbox will change the color of all his UI_Objects when the user hover over it and again when the user enters the activehand state while hovering, such as closing his hand. In the former situation the activation callback function is also called. No instance of this object is currently used in the interface.

The Abstr_HitboxSlideButton class is a different abstract class that adds behavior to the parent class. Specifically it adds functions to calculate how far the hand that is interacting with the hitbox has moved since the last frame. This information can be used in its children to move hitboxes and UI_Objects relative to the hand's movement. Parameters within this class define how far the hand can move from the hitbox's original position and when the activation callback function is called.

The UI_HitboxScrolbar class adds a vector of Abstr_UI_Hitboxes and uses the behavior defined in UI_HitboxSlideButton to move them up or down according the movement of interacting hand.

COMMENT: redefinition of the draw method

The UI_HitboxLockScrolbar class adds the act of locking each Abstr_UI_Hitboxes derived object in its vector into predefined positions when the user stops scrolling. It also manage the behavior of the orange selection box as discussed in the previous chapter where this class of hitbox is used as the scrollbar in the end design.

The `UI.HitboxSlideButton` defines behavior to slide a button until an activation point is reached. In this specific case the user has to put his hand into the `handActive` state, before he is able to slide the button. The rope described in the chapter **Reference!** is an instance of this class.

The `UI.HitboxHoverSlideButton` basically does the same as the `UI.HitboxSlideButton` but doesn't require the user to put his hand into the `handActive` state. A hand hovering in any state on the right side of the hitbox rectangle will begin sliding it towards the activation point. This class is used in the chapter **Reference!** as the delete slide button for the recorded elements in the scrollbar.

The `Abstr.Screen` class is an abstract class for the creation of different screens. It holds a vector of shared pointers to all hitboxes on the screen, defines how the puppet of the user is drawn and provides functions to draw background and top UI elements to be defined by his children. For every different screen a new child of this abstract class needs to be added.

The `RecordScreen` class is a child of `Abstr.Screen` that defines the type and the properties of the hitboxes and `UI.hitboxes` that are used to create the screen as it is described in the previous chapter. The hitboxes are also given their callback functions that are used when the hitbox is activated and when the Model calls for an update of the UI representation.

4.1.2 Flow of the program

The general flow of the program is that at the start of the program, after the creation of the UI and Model objects, the UI class object is instructed to fill his pointer to the `Abstr.Screen` class with an instance the `RecordScreen` class and then call the function `createScreen()` on it to create all the hitboxes and UI-objects that belong to that particular screen. Then the program loop is started and the `processBody()` method in the Model class calls the `DrawFrames(etc)` method in UI which will pass it on to its pointer of `Abstr.Screen` now pointing to a `RecordScreen` object. This will first call the `draw()` method on all the hitboxes belonging to the background, which will recursively pass on the `draw()` function to any `Abstr.UI.Hitbox` child or `UI.Object` child it contains. At this stage only the `UI.Object` class and his children will actually call the `D2D.Graphics` object to draw something on the `Direct2D` canvas. Then `drawFrames(etc)` function continues by calling the `D2D.Graphics` object to draw the puppet representing the user and on top of that the hitboxes of the top UI are drawn in the same manner as the background. Last the coordinates and states of the hands of the user are given to every hitbox in the screen, these will use this information to determine and execute any changes in their state or their UI-objects based on their criteria. If a hitbox determines that his activation criteria have been met he calls his `activateFunctionCallback(etc)` that will make calls directly to the Model. This happens every cycle when the `processBody()` function in the Model class is called, which is usually 30 times per second.

Whenever something changes in the Model structure that affects the representation in the UI, a flag is set and during `processBody()` the `updateHitboxes()` function is called on the UI object before the UI is drawn. The UI passes this on the `Abstr.Screen` child who calls an `update()` function on all the hitboxes it contains. Each hitbox will call their own `updateFunctionCallback(etc)` to retrieve all information related to them from the Model and process it. Hitboxes that contain other hitboxes will call the `update()` function recursively.

4.2 Back-end software

4.2.1 Class diagram

The focus of the application is reflected by the structure of the code. The class diagram shows a model-view pattern to make a distinction between the graphical interface of the application and the back-end (see figure 4.3).

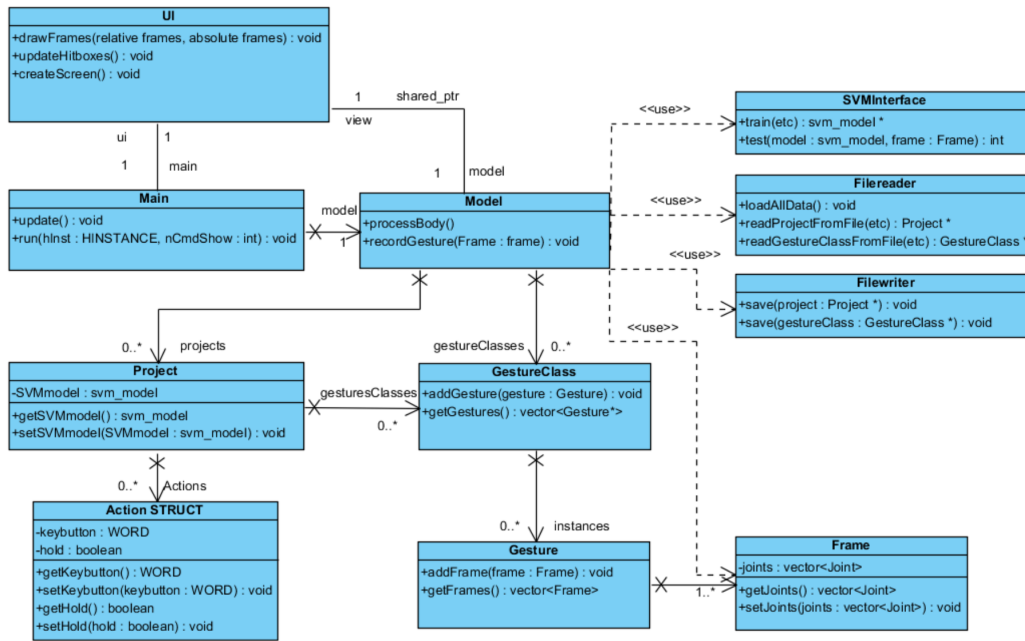


Figure 4.3: The class diagram of the application, focusing on the back-end

The Kinect camera has the ability to identify and measure the position relative to the camera of 25 points of a person, for instance: left elbow, right elbow, head, center of the spine, left wrist, . . . All of these points are referred to as `Joints` and can be accessed using the libraries that come as part of the Kinect SDK installation. Each `Joint` has an x , y and z coordinate, so it is unambiguously defined in space.

When the Kinect measures all of the `Joints` at one specific point in time, these 25 `Joints` together form one frame. This can be seen as a single picture of a person taken by the camera. Analogous to a movie, which is nothing more than a quick succession of pictures, an instance of `Gesture` collects all `Frames`, which all together, contain information about how the person moved over a period of time.

In order to improve the application's ability to recognize an exercise, each of the exercises must be trained more than once by the physical therapist. For each exercise that is trained, an instance of `Gesture` is created. The `Gestures` that contain data about different executions of the same exercise are grouped into a `GestureClass`. In other words, one `GestureClass` object contains all trainings of the same exercise.

After setting up a `Project` object, it contains all data that is needed for playing a game using the application. It has a map that maps a label for each exercise to a `GestureClass` and the actions linked to that `GestureClass`. An `Action` contains information about the keyboard key that needs to be pressed and if that key should be held down or be quickly pressed and released. The `Model` class forms the core of the application. It controls the flow of the program and contains all important objects, such as the `Project` and the `GestureClasses`. By storing the `GestureClass` objects in `Model` rather than `Project`, it is possible to reuse the same `GestureClasses` in multiple `Projects`.

`SVMInterface` takes all gesture data and converts it to a format that is accepted by the LibSVM library [14]. This is an existing library that provides a C++ support vector machine implementation and is used in this application for gesture recognition.

To prevent having to train all exercises again each time the application is started, all necessary data is saved into files in the data folder using the `Filewriter` implementation. This includes the data of the

`Gestures`, `GestureClasses`, `Project` and the computed SVM model. The `FileReader` is responsible for reading data from the saved data files when the application is started.

4.2.2 Flow of the program

After initialization of variables like the instances responsible for the user interface and the communication with the Kinect camera, the program locks into an infinite loop while the application is running. This main loop consists of three processes: fetching new data from the Kinect camera, analyzing this data in order to use it for recording or predicting gestures and updating the GUI with relevant changes.

It is important that none of these three processes block the continuous flow of the main loop. The update rate of the interface is only as fast as the rate with which the main loop is executed. If it is slowed down with a blocking or long-running loop, the GUI appears to stutter or freeze and no new data is collected from the Kinect camera.

The Kinect can only provide new data as fast as 30 times per second. If no new data is available at the moment the main loop requests the data, for instance when the main loop is executed faster than 30 times per second, the program skips this process and continues with the other two processes.

To control the flow of the program without reverting to blocking loops, software flags are used. These flags keep track of the current state of the program. By evaluating the state of the flags, it is possible to indicate if the program is predicting a gesture, recording a gesture or animating specific graphical elements of the GUI.

4.3 Gesture recognition

It is possible to split up the implementation of gesture recognition into two parts: the recording of a gesture by the physical therapist and the prediction of a gesture executed by a patient.

4.3.1 Recording a gesture

Code snippet 4.1 shows the code for recording a gesture. When the recording of a gesture is initialized, the first frame is stored for reference. Actual recording starts when the users starts moving. This means that the first frame is different from the current frame. Two frames are considered equal when all of the joints of both frames in all three dimensions have a difference of less than 12 centimeters. In other words, if the person moved more than that, the frames are not equal. The threshold of 12 centimeters is chosen due limitations of the Kinect camera, where small variations in measurements are possible. Additionally, it is hard for a person to stay perfectly still, so the choice of the threshold can eliminate recognizing small tremblings as the user having moved.

When actual recording is initialized, the frames are stored in a buffer. Together, the frames of this buffer form an entire gesture. If the buffer is filled up without the user moving, this means that the user is recording a posture rather than a gesture. However, when the user moves, recording can be ended by standing still in an ending position of choice for approximately 1 second.

```
void Model::recordGesture(Frame & frame) {
    if (!initialized) {
        //Set the first frame on the first call
        firstFrame.setFrame(frame);
    }
}
```

```

    initialized = true;
}

if (!startedMoving) {
    if (! frame.equals(firstFrame)) {

        startedMoving = true;
        framesBuffer.clear();
    }
    else if (framesBuffer.size() > NOT_MOVING.FRAME_DELAY) {
        addRecordedGesture();
    }
    else {
        return;
    }
}

//Add the frame to the buffer
framesBuffer.push_back(frame);

if (framesBuffer.size() > NOT_MOVING.FRAME_DELAY &&
    framesBuffer.back().equals(framesBuffer.at(framesBuffer.size() -
        NOT_MOVING.FRAME_DELAY))) {
    addRecordedGesture();
}
}

```

Code snippet 4.1: method to record a gesture

4.3.2 Predicting a gesture

Code snippet 4.2 shows a method for predicting a gesture. This method is called recursively. Gestures are split up into a number of smaller parts, referred to in the code as `NB_OF_LABEL_DIVISIONS`. If this constant equals four, it means that a gesture is split up into four parts. For instance, if the base label of a gesture equals one, predictions for labels one, two, three and four all relate to one of the parts of the gesture.

Continuing with this example, if during prediction a part of a gesture with label four is found, the recursive method `isGestureExecuted` is called. It checks if shortly before label four is found, label three is predicted as well. If label three is predicted, the method is called recursively to check for label two. This goes on until the base label is reached. At this point, the recursion is ended and the gesture is recognized as being successfully executed. All parts of the gesture must be executed in the right order and fast enough. If the gesture is executed too slow, the buffer with predicted labels is cleared.

The method allows that one label is not predicted correctly. There are many reasons this can happen, such as noisy measurements, the user moving differently or not having trained the gesture sufficiently. In the case of the considered example, if labels one, three and four are predicted, and label two is missing, the gesture is still considered to be executed successfully. This can improve user experience while playing a game, as a gesture has less of a chance of being not recognized wrongfully.

```

bool Model::isGestureExecuted(int checkLabel, int posInBuffer, int recursiveCounter, int
    badCounter) {
    //The label exists and is linked to a posture.
    if (activeProject->containsLabel(checkLabel) &&
        activeProject->getGestureClass(checkLabel)->getGestures().front()->isPosture())
        return true;
}

```

```
//Done enough recursive checks to confirm the gesture has been executed.
if (recursiveCounter >= NB_OF_LABEL_DIVISIONS)
    return true;

int nextLabelToCheck = checkLabel - 1;
for (int i = posInBuffer; i >= 0; i--) {
    if (labelsBuffer.at(i) == nextLabelToCheck)
        return isGestureExecuted(nextLabelToCheck, i, recursiveCounter++, badCounter);
}

//If this point is reached, a label that is one less than the given
// label cannot be found in the buffer.
//The gesture may still have been executed, so keep checking for the
// next one if the badCounter is not too high.
if (badCounter > 0)
    return false;

return isGestureExecuted(nextLabelToCheck, posInBuffer, recursiveCounter++, badCounter++);
}
```

Code snippet 4.2: method to verify if a gesture with given label is executed

Chapter 5

Results

Usability tests make it possible to verify if the design meets the expectations of the target group and offer a way to discover unexpected problems that are hard to track by the developers themselves. Test persons have a fresh look on the application and can provide useful insights in the way they interpret visual elements of the user interface.

In addition, the technical evaluation considers the effect of the implementation, ensuring that the user is not hindered by performance issues or similar problems.

5.1 Usability test

The graphical user interface is designed for use by physical therapists. As such, acquiring feedback from therapists is essential in developing an interface that is both useful and easy to use. The process of the usability test is described, followed by an overview of the most important results of the evaluation itself and an interpretation of these results.

5.1.1 Description of the test

A usability test is conducted with physical therapist Dries Lamberts at Windekind Leuven, an organization that focuses on the education of children with disabilities and offers a full program to assist them with their specific difficulties. For this test, the Kinect camera is connected to a computer that runs the application and, positioned on a desk at hip height. The user indicates that he is familiar with using the Kinect camera and motion-based games.

A short introduction on the purpose of the application is given to the test person for reference. After that, he is presented with several tasks and is asked to apply the think-aloud protocol in order to obtain as much information as possible about the interaction with the application and the thoughts of the user while doing so.

In preparation of the usability test, gestures are pre-recorded that both serve the purpose of testing the robustness of the application as well as having a fallback plan in case problems arise during the test.

One task is to watch the pre-recorded gestures and control a game called Space Invaders using the pre-recorded gestures. This is a game that has an avatar move left or right on the screen and shoot bullets (see figure 5.1). As such, four gestures are pre-recorded: one gesture that allows the avatar to move to the left, one to move to the right, one to shoot bullets and a neutral gesture that is linked to no keyboard key. At this point, it is explained to the user why a neutral gesture is required, without going into detail

about the specifics of SVM. The goal of this task is twofold. Firstly, the task is used to verify if the test person is able to watch replays of the pre-recorded gestures without any guidance and to learn what gestures are pre-recorded without additional information. After this is done, it is explained to the user how the gestures and keyboard keys are linked together. Secondly, it makes it possible to confirm if a correct prediction of a gesture execution does not depend on the person performing the gesture, as problems can arise due to the person recording the gestures and the person playing a game having a very different physique.



Figure 5.1: Screenshot from *Space Invaders*

Another task is to play the game Sokoban Geek. This puts the gamer in control of an on-screen avatar that can move up, down, left or right. Walking into boulders allows the player to move them in order to solve a puzzle (see figure 5.2). The game is shown to the test person and after trying the game, he decides how many gestures are required to play the game. Then, he comes up with different gestures for each input and records three trainings for each of them. The test person is asked to delete a recorded gesture and record it again. After recording all gestures, he tries to play the game. The goal of this task is to check how the test person reacts to the recording elements of the interface and if it is clear how to interact with them without assistance.




Figure 5.2: Screenshot from *Sokoban Geek*

With permission of the test person, during the entire test, the computer screen is recorded, as well as

everything said, for analysis purposes. The system usability score (SUS) allows to obtain a numeral indication of the usability of the application based on ten questions. At the end of the usability test, the user is asked to fill out a questionnaire with these ten questions.

5.1.2 Evaluation

For the first task, the user notices the recorded gestures in a list on the right-hand side of the screen. By trying to move its avatar's hand to align with the previews of the recordings, he notices that it is possible to scroll through the list of recordings. He also notices the delete option next to the recordings list and indicates that he is anxious about deleting one of the recorded gestures.

Watching a replay of the recorded gestures initially is not clear to the user. He tries to drag one of the gestures to the square on the left hand-side of the screen. He notices that this is not possible, but makes several attempts do this. After that, he indicates that he can't find it and looks at the interface without trying to interact with anything. To replay a recorded gesture, the user can point with the avatar's hand to the recorded gesture in the green area of the scrollbar. At that point, a replay of that gesture is displayed automatically. The replay stops while scrolling through the list or when not pointing at the gesture in the green area. The user did this several times during this task subconsciously, without noticing that a replay started playing at the left side of the screen. The user required assistance to complete this task. He was able to do this after knowing that the gesture in the green area is displayed at the left-hand side square.

After watching replays of the recorded gestures, the user is able to replicate in real life what gestures are recorded and displayed on-screen. When the Space Invaders game is started, the user has no problem interacting with the game and does not require any assistance. He is able to control the in-game avatar subtly and meet the game objectives. After asking for his experience with the controls, he indicates that the controls are responsive and that he does not experience any lag.

For the second task, the user gets the chance to see what the Sokoban Geek game looks like and how it is played. After that, he states that he needs five gestures to play the game. Four gestures are linked to a directional keyboard key and one gesture that serves as the neutral gesture. To record a gesture, it is required that the user points the hand of the on-screen avatar to the grip of the pulling cord, makes a fist, and move his fist down to start recording. This translates to pulling the displayed cord. Trying to start recording the first gesture, the user pulls the cord without any problems. The recording screen appears and he starts performing a gesture when the interface indicates that recording has started. He stands still when he is done with the gesture and the recording stops automatically.

Upon being asked to delete the gesture he just recorded, he needs no assistance to push the gesture object to the right into a box indicating that the gesture can be deleted. He indicates that his earlier anxiousness of deleting a gesture by accident is unnecessary, as it requires some effort to perform this deleting action and it is hard to do it by accident.

When trying to record a gesture the second time, the user is unaware that he is required to close his hand into a fist to pull the cord. He tries pointing at the grip and moving his hand down and indicates that he isn't sure why recording does not start this time. Only after several tries, he tries closing his hand to pull the cord, activating the recording screen. He states that he did not know that closing the hand was required and that he did that by accident to record the first gesture. He indicates that he has experience using the Kinect camera for other games and that these games never required to close a hand, so he was not sure that the Kinect could make the distinction between open and closed hand.

The user notices clearly that the recording of a gesture stops when he stands still. Because he was not

aware of this, one gesture was not recorded the intended way, but the user was able to delete this gesture and record a new one without assistance.

When trying to watch a replay of a recorded gesture, he is not sure if he needs to point at the gesture with open or closed hand. Although both cases function in the same way, the user states that using a closed hand feels like the application responds better.

After the user records all gestures, he is able to play the game. However, the gesture that is linked to the left arrow key does not respond. The user is asking if the key is linked correctly to the gesture and thinks this is the problem. He misses feedback about what gesture is linked to what key. After verifying for the user that the key is linked correctly, we present him with the question of how he would handle this problem without any assistance. He suggested to delete the recorded gestures and record them again. After doing this, the user is able to play the game without further problems.

The entire user test approximately took 45 minutes. After the test, the user asked because of his own experience with disabled children if very small gestures are supported by the application and if the application can be used for other purposes than playing games. An example he thought of was using gestures to type words using sign language. He also added that he was glad that use of the space bar is supported by the application, as he has experience with similar projects that do not support the space bar, while it is one of the most used keys in computer games. However, a lot of games the children are interested in require the ability to point with a cursor on-screen. This is a feature he would appreciate. Finally, he stated that he is very interested in using the application, allowing disabled children to play games and exercise while doing so.

5.1.3 Interpretation

The physical therapist proved to understand the concept of the application. This shows that it is possible for persons without programming knowledge to use the application for recording gestures and controlling computer games. He also understands that the use of the application is not limited to games and can also be used for other computer applications that require keyboard input.

A few tries are required to get familiar with how to start recording and how to end it. In particular, without prior knowledge or experience, more feedforward is required to communicate that the hand needs to be closed when pulling the cord and that the user needs to stand still to stop recording. However, after a few tries, it is possible to record multiple gestures in succession without any of them not being recorded as intended.

Replaying recorded gestures requires the user to align his on-screen hand with the gesture inside the green area of the scrollbar. There is no form of feedforward indicating that performing this action results in the gesture being replayed. The feature is intended for the user to be discovered while scrolling through the gestures or while trying to delete a gesture. As such, this can be confusing to the user.

Both the actions of recording gestures and replaying recorded gestures are performed faster after having experienced this firsthand. The provided feedforward is too subtle to immediately make it clear what actions are expected. To improve this, it is either possible to provide more explicit feedforward, or to introduce the user to the most important functions using a tutorial.

5.2 Technical evaluation

The application is developed with the Microsoft Visual Studio integrated development environment (IDE) in C++ and can run on Windows operating systems. This limitation is due to the use of the IDE's visual interface editor and Direct2D, as well as the feature of saving and loading data files. The application is tested to run on both Windows 8.1 and Windows 10 machines.

All of the required files for running the application, including resource images, take up approximately 20 MB of space. Saving all exercises needed for playing a game takes up approximately an additional 0.6 MB of space. This depends on the number of exercises recorded, the number of trainings done for each exercise and the duration of each exercise.

The application refreshes the interface at a rate of 30 frames per second (fps), which is limited due to the use of the Kinect camera, sampling 30 times per second. The Kinect also introduces limitations related to the environment and the user. These limitations are that overly lit rooms make it harder for the camera to track a person. The same is true for persons wearing black or reflecting clothing. These findings are in accordance to Microsoft's guidelines for using the Kinect [15].

Chapter 6

Discussion

6.1 Reflection on the results

Usability tests are import tools for the evaluation of a prototype. However, with a limited number of test persons, a quantitative analysis is not possible due to the necessity of a test group that is sufficiently large to produce statistically useful data. Even for a qualitative analysis, using the opinion of only one physical therapist may result in responses that are biased towards a certain view the therapist has on the subject. To distinguish between genuine design defects and problems only encountered due to individual bias, it is necessary to have a large enough test group. For the scope of this thesis, the number of user tests is limited, so a more careful approach is needed when interpreting the results of the test.

The test users, including the physical therapist, require time to get familiar with non-traditional gesture-based interface elements. The time needed depends on the person and his knowledge and experiences with other gesture-based interfaces. An important part of this is knowledge about the Kinect camera and its possibilities. The test users are worried of doing something that the application does not know how to process. A common problem is that it is unknown to the users that the Kinect can make a distinction between an open and a closed hand, even while there is visual feedback provided when opening or closing a hand: the hands on-screen turn respectively green or red. This means that the user does not consider opening or closing his hands when trying to interact with something on the screen. This problem is always related to the features supported by the used sensors. The Kinect camera supports this feature, but similar cameras may not. Because of this, it is not possible to simply assume user knowledge of this operation.

When test persons get the time to experiment with the interface at their own pace, they eventually find a way by themselves to interact with the application. Due to the limited amount of different interactive elements, the application has a gentle learning curve. A tutorial can help bringing new users up to speed with how to use the application. However, a tutorial can also be a source of frustration or confusion for the user if not done well. As such, tutorials should be seen more like a last resort instead of the go-to solution for any interface-related problem. More and better feedback and feedforward is generally the preferred option to choose, as it increases natural coupling between the required actions and its function [5]. Also, providing more than one form of feedforward for the same action can help making the interaction more clear as well. For instance, the color of an element changes when the users hovers over that element, implying that interaction is possible. This can be expanded by having a sound effect play at the same time as the on-screen color change.

A potential problem is the limited amount of space the user has in the room he is using the application. Unless the user is required to move left or right during the use of the application, all interactive elements must be placed within reach of the user's avatar. This means that non-interactive elements can be placed further away from the avatar. This introduces the problem of limited interface space on the screen and

can result in elements that are functionally coupled together to be separated. For instance, the scrollbar with recorded gestures is on the right of the user's avatar, but the playback of the gesture is shown on the left. It limits the locational aspect of the functional feedback, but at the same time indicates that no specific action is required for interacting with the gesture display on the left as it is out of reach. For instance, to some extent, this can be compared to having to interact with a Blu-ray player, but receiving visual feedback from another place, i.e. the television screen.

Unrelated to the interface design, but an important part of the application is gesture recognition. Due to the chosen approach, it is possible to record and train exercises as gestures as well as postures. However, the way SVM works limits the requirement for high accuracy when executing a gesture or a posture. The executed exercise is always recognized as the recorded exercise that looks the most similar. If all of the recorded exercises are very dissimilar, exercises do not need to be accurately executed in order for them to be recognized as a specific exercise. This means that it is not possible to tell if an exercise is executed correctly. It requires the physical therapist to keep a close look at the way the patient performs an exercise during play.

Furthermore, the speed with which an exercise is executed is not checked. Only if the execution is too slow, it is not recognized as a recorded exercise. Faster execution does not affect the gesture recognition, on condition that the Kinect's sampling limitations are not surpassed. To solve this problem, as well as the previous problem of inaccurate executions, the therapist needs to record not only the required exercises, but also examples of how exercise should not be executed, i.e. exercises that are performed too slow, too fast or just not the way the gesture has to be done. This requires the therapist to record many examples of wrong executions, which requires a lot of time to do and eventually leads to a less pleasant experience during play as far fewer executions are accepted as being correctly executed. In addition, it requires the therapist to have some understanding of how machine learning works, which does not match the goal of designing an application that can be used without any programming knowledge.

6.2 Reflection on the process

There are different approaches for designing an interface, but it always comes down to not treating the interface as an afterthought. As such, integrating the interface design right from the start with back-end software design is essential. For instance, requirements concerning the responsiveness of the interface are for the biggest part set by the way the back-end software is designed and implemented.

However, it is necessary to switch fast from interface prototyping to an actual implementation. Since the early prototypes are static and non-interactive, it is only possible to really evaluate a decision after having an implementation of the concept and to obtain feedback from the users it is intended for. In other words, while static prototypes can undergo several iterations until they are *perfected* in theory, only implementations in practice can really convey the feel of the interface.

To come up with ways of interacting with the application other than relying on traditional controls, an experimental approach is tested. The idea is to interview a physical therapist in order to obtain ideas from him concerning the interaction with the application. This makes it possible to find solutions that fit the needs and expectations of the user. Early tryouts of this experimental technique show that persons are reluctant of coming up with ideas they are not familiar with. It is more assuring and easier to fall back on traditional, well-known elements.

In the same sense, user testing often leads to results implying that more conventional interactive elements are easier and quicker to understand than others, even if they rely on making use of the mime pattern

discussed before. This can be related to years of experience with other types of interfaces and confirms that there is need of a more streamlined way for gesture-based interaction that unifies all interfaces of this kind [3].

Including a scrollbar in a gesture-based interface is an example of this. The user has experience using scrollbars, so there are no problems to recognize the visual element as a list that can be scrolled through, but interacting with it is more difficult initially. This can be related to differences in the way gesture-based and touchscreen-based interfaces work. On touchscreens, it is possible to swipe with a finger on the screen to scroll, then lift it and touch the screen again to scroll again. With gesture-based interfaces, the action analogous to lifting the finger is less obvious to achieve. One way is to have the user move his hand out of the scrollbar before he can scroll in the same direction again. This is how it is implemented for the application. A different approach is to use composite gestures. The user can close his hand to grab the scrollbar and scroll through it, then open his hand, reposition it and close it to scroll again. This allows the action of scrolling to be faster, but it also adds the burden of having to remember more actions for interacting with this element and it can be tiresome to close and reopen the hand many times in succession. In the end, this comes down to personal preference and requires more user tests before a more general conclusion can be made.

6.3 Future work

Physical therapist Dries Lamberts indicated that the action of recording and training gestures can be pleasant for patients to do it themselves. This is especially useful if the developed application is not exclusively used for rehabilitation, but also as a way to have some physical exercise on a regular basis, as is the case for children with physical disabilities. As the GUI of the application is focused on being used by therapists, a different interface is required to simplify and optimize its use for children. This includes features like pressing a keyboard key to be designed appropriately.

A second possibility for future work is to use the current back-end structure to allow the user to create patient-specific projects, to which all exercises are coupled. On startup of the application, it can be possible to require the user to enter the name of a patient to load the project with all exercises they used before. This decreases the time needed before a patient can start exercising and playing a game.

Chapter 7

Conclusion

The developed application succeeds in offering a way to incorporate exercises with playing games. A physical therapist can come up with exercises that fit the needs of the patient and link each of these exercises to a keyboard key. As such, any game that can be played with keyboard controls can be played with this application, independent of the chosen exercises. This gives patients a very large pool of games to choose from and helps keeping them motivated to do all needed exercises.

The therapist can train all required exercises without any programming knowledge. Machine learning algorithms are used to evaluate which of the exercises is performed by the patient and allow the therapist to support the patient in a different way, by focusing together with the patient on playing the game instead of counting the number of times an exercise is done. The way SVM is used to recognize the executed exercise minimizes the time between a gesture execution and a reaction from the application. This allows for a more pleasant experience for the patient while playing games, as they feel more responsive.

Special attention is paid to the application running smooth during its use. The application updates the user's avatar drawn on the interface at a rate of 30 fps, which is as fast as the Kinect 2.0 camera can measure new data. As an important form of feedback, this helps mimicking user movement in real-time and keeps the user immersed into the experience.

The GUI uses visual elements that express their function through feedforward. The goal is to use objects that users can recognize from real life and inherently know how to interact with. A gentle learning curve paves the way for new users in minimizing the amount of additional information needed before the application can be used. By naturally coupling action to reaction, feedback informs the user of how an interaction affects the application and makes it easier to understand the result of it.

Bibliography

- [1] Brauner P, Valdez AC, Schroeder U, Ziefle M. Increase Physical Fitness and Create Health Awareness through Exergames and Gamification. *Human Factors in Computing and Informatics*. 2013;7946(1):349 – 362.
- [2] Jacob RJK, Girouard A, Hirshfield LM, Horn MS, Shaer O, Solovey ET, et al. Reality-based interaction: a framework for post-WIMP interfaces. *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. 2008;CHI 08(1):201 – 210.
- [3] Norman DA. Natural User Interfaces Are Not Natural. *Interactions*. 2010;17(3):6 – 10.
- [4] Shneiderman B. Direct Manipulation: A Step Beyond Programming Languages. *IEEE Computer*. 2010;16(8):57 – 69.
- [5] Wensveen SAG, Djajadiningrat JP, Overbeeke CJ. Interaction frogger: a design framework to couple action and function through feedback and feedforward. *Proceedings of the 5th conference on Designing interactive systems*. 2004;DIS 04(1):177 – 184.
- [6] Hondori HM, Khademi M. A Review on Technical and Clinical Impact of Microsoft Kinect on Physical Therapy and Rehabilitation. *Journal of Medical Engineering*. 2014;2014(1):16.
- [7] Lange B, Koenig S, McConnell E, Chang CY, Juang R, Suma E, et al. Interactive game-based rehabilitation using the Microsoft Kinect. *Virtual Reality Short Papers and Posters*. 2012;p. 171 – 172.
- [8] Chang YJ, Chen SF, Huang JD. A Kinect-based system for physical rehabilitation: A pilot study for young adults with motor disabilities. *Research in Developmental Disabilities*. 2011;32(1):2566 – 2570.
- [9] Peng W, Lin JH, Crouse J. Is Playing Exergames Really Exercising? A Meta-Analysis of Energy Expenditure in Active Video Games. *Cyberpsychology, Behavior, and Social Networking*. 2011;14(11):681 – 688.
- [10] Staiano AE, Calvert SL. Exergames for Physical Education Courses: Physical, Social, and Cognitive Benefits. *Child Dev Perspect*. 2011;5(2):93 – 98.
- [11] Dahl-Popolizio S, Loman J, Cordes CC. Comparing Outcomes of Kinect Videogame-Based Occupational/Physical Therapy Versus Usual Care. *Games Health J*. 2014;3(3):157 – 161.
- [12] Annema JH, Verstraete M, Abeele VV, Desmet S, Geerts D. Video games in therapy: a therapists's perspective. *International Journal of Arts and Technology*. 2013;6(1):106 – 122.
- [13] Microsoft. Kinect for Windows SDK 2.0. Microsoft; 2014. <https://www.microsoft.com/en-us/download/details.aspx?id=44561>.
- [14] Chang CC, Lin CJ. LIBSVM - A Library for Support Vector Machines; 2016. <https://www.csie.ntu.edu.tw/~cjlin/libsvm/>.

- [15] Microsoft. Human Interface Guidelines v2.0. Microsoft; 2014. <http://download.microsoft.com/download/6/7/6/676611B4-1982-47A4-A42E-4CF84E1095A8/KinectHIG.2.0.pdf>.

FACULTY OF ENGINEERING TECHNOLOGY
CAMPUS GROUP T LEUVEN
Andreas Vesaliusstraat 13
3000 LEUVEN, Belgium
tel. + 32 16 30 10 30
fax + 32 16 30 10 40
fet.groupt @kuleuven.be
www.fet.kuleuven.be



MEMBER OF **ASSOCIATIE
KU LEUVEN**